

Arth Task 7.1 : Elasticity Task

: Increase the storage of a data-node to increase the overall storage: Integrating LVM with Hadoop



First we create ah Hadoop Cluster with 1 master and 1 Datanode

```
Datanodes available: 1 (1 total, 0 dead)

Name: 65.0.71.177:50010
Decommission Status : Normal
Configured Capacity: 5217320960 (4.86 GB)
DFS Used: 8192 (8 KB)
Non DFS Used: 306167808 (291.98 MB)
DFS Remaining: 4911144960 (4.57 GB)
DFS Used%: 0%
DFS Remaining%: 94.13%
Last contact: Sat Mar 13 07:22:17 UTC 2021
```

Before Increasing or decreasing the memory , we check the total number of partitions in the OS using : **fdisk -l**

```
[root@ip-172-31-13-235 ec2-user]# fdisk -l
Disk /dev/xvda: 10 GiB, 10737418240 bytes, 20971520 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: 246B752E-8CB4-41E7-B9B1-365A93ACF890

Device        Start      End  Sectors  Size Type
/dev/xvda1    2048      4095     2048    1M BIOS boot
/dev/xvda2   4096 20971486 20967391   10G Linux filesystem
[root@ip-172-31-13-235 ec2-user]#
```

Next we add a Secondary storage to the OS : Since am using AWS , I use EBS to add a secondary storage

Attach Volume



Volume ⓘ vol-0eeba986d16a3089c in ap-south-1b
Instance ⓘ in ap-south-1b
Device ⓘ
Linux Devices: /dev/sdf through /dev/sdp

Note: Newer Linux kernels may rename your devices to /dev/xvdf through /dev/xvdp internally, even when the device name entered here (and shown in the details) is /dev/sdf through /dev/sdp.

Cancel Attach

```
[root@MiWiFi-R3L-srv ~]# fdisk -l
Disk /dev/sda: 20 GiB, 21523628032 bytes, 42038336 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x8628e552

Device      Boot    Start        End    Sectors    Size Id Type
/dev/sda1   *            2048    2099199    2097152     1G 83 Linux
/dev/sda2                2099200  42037247  39938048    19G 8e Linux LVM

Disk /dev/mapper/rhel-root: 17 GiB, 18291359744 bytes, 35725312 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes

Disk /dev/mapper/rhel-swap: 2 GiB, 2155872256 bytes, 4210688 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
[root@MiWiFi-R3L-srv ~]#
```

Now we have to create a physical volume for the storage using command: **pvccreate /dev/xvdf** and to confirm we use command : **pvddisplay**

```
[root@ip-172-31-13-235 hadoop]# pvcreate /dev/xvdf
Physical volume "/dev/xvdf" successfully created.
[root@ip-172-31-13-235 hadoop]# pvdisplay
"/dev/xvdf" is a new physical volume of "10.00 GiB"
--- NEW Physical volume ---
PV Name                /dev/xvdf
VG Name
PV Size                10.00 GiB
Allocatable            NO
PE Size                0
Total PE               0
Free PE                0
Allocated PE           0
PV UUID                vYZ5lb-bma1-nMKq-lLiJ-t6Zd-u44Q-mbirPJ
```

After creating Physical volume , we need to create a volume group using command : **vgcreate [Group-name] /dev/xvdf** and to confirm we use command :**vgdisplay**

```
[root@ip-172-31-13-235 hadoop]# vgcreate Rahul /dev/xvdf
Volume group "Rahul" successfully created
[root@ip-172-31-13-235 hadoop]# vgdisplay
--- Volume group ---
VG Name                Rahul
System ID
Format                lvm2
Metadata Areas         1
Metadata Sequence No   1
VG Access              read/write
VG Status              resizable
MAX LV                 0
Cur LV                0
Open LV                0
Max PV                 0
Cur PV                1
Act PV                 1
VG Size                <10.00 GiB
PE Size                4.00 MiB
Total PE               2559
Alloc PE / Size        0 / 0
Free PE / Size         2559 / <10.00 GiB
VG UUID                WWfPdM-nLx0-P6kd-d5hN-pTLh-mwbU-hAnFzJ
```

Finally creating a Logical Volume using the previous created volume group using command : **lvcreate — size [Storage amount] — name [group-name] [physical volume]**

```
[root@ip-172-31-13-235 hadoop]# lvcreate --size 5G --name LVM1 Rahul
Logical volume "LVM1" created.
[root@ip-172-31-13-235 hadoop]# lvdisplay
--- Logical volume ---
LV Path                /dev/Rahul/LVM1
LV Name                 LVM1
VG Name                 Rahul
LV UUID                 CD8cud-wuf7-uwCp-09YL-EURE-tbKh-GcCIqd
LV Write Access         read/write
LV Creation host, time ip-172-31-13-235.ap-south-1.compute.internal,
3 07:18:37 +0000
LV Status                available
# open                  0
LV Size                 5.00 GiB
Current LE              1280
Segments                1
Allocation              inherit
Read ahead sectors      auto
- currently set to      8192
Block device            253:0
```

We need to format the volume before mounting using command :

mkfs.ext4 /dev/[Groupname]/[partition name]

```
Disk /dev/mapper/Rahul-LVM1: 5 GiB, 5368709120 bytes, 10485760 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
[root@ip-172-31-13-235 hadoop]# mkfs.ext4 /dev/Rahul/LVM1
mke2fs 1.45.6 (20-Mar-2020)
Creating filesystem with 1310720 4k blocks and 327680 inodes
Filesystem UUID: ea33cfda-0f30-4429-8f97-230429306fef
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736

Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done
```


Finally , we mount the volume using command :

mount /dev/[Partition-name]/[Group-name] /storage-name

```
[root@ip-172-31-13-235 etc]# mount /dev/Rahul/LVM1 /dn
[root@ip-172-31-13-235 etc]# df -hT
```

Filesystem	Type	Size	Used	Avail	Use%	Mounted on
devtmpfs	devtmpfs	378M	0	378M	0%	/dev
tmpfs	tmpfs	403M	0	403M	0%	/dev/shm
tmpfs	tmpfs	403M	16M	388M	4%	/run
tmpfs	tmpfs	403M	0	403M	0%	/sys/fs/cgroup
/dev/xvda2	xfs	10G	1.8G	8.3G	18%	/
tmpfs	tmpfs	81M	0	81M	0%	/run/user/1000
/dev/mapper/Rahul-LVM1	ext4	4.9G	20M	4.6G	1%	/dn

Thus the storage has been increased , to confirm we can use command :

hadoop dfsadmin -report

```
[root@ip-172-31-14-252 hadoop]# hadoop dfsadmin -report
Configured Capacity: 10724814848 (9.99 GB)
Present Capacity: 8818454528 (8.21 GB)
DFS Remaining: 8818446336 (8.21 GB)
DFS Used: 8192 (8 KB)
DFS Used%: 0%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0

-----
Datanodes available: 1 (1 total, 0 dead)

Name: 65.0.71.177:50010
Decommission Status : Normal
Configured Capacity: 10724814848 (9.99 GB)
DFS Used: 8192 (8 KB)
Non DFS Used: 1906360320 (1.78 GB)
DFS Remaining: 8818446336 (8.21 GB)
DFS Used%: 0%
DFS Remaining%: 82.22%
Last contact: Sat Mar 13 06:50:46 UTC 2021
```

Thus it is confirmed that the storage has been increased by 5 GB.

This whole process can be automated using python scripting :

GITHUB URL : [https://github.com/Rahul1498/ARTH-
Tasks/blob/main/LVMautomation.py](https://github.com/Rahul1498/ARTH-Tasks/blob/main/LVMautomation.py)